

# Big Data Meets HPC Log Analytics: Scalable Approach to Understanding Systems at Extreme Scale

Byung H. Park\*, Saurabh Hukerikar\*, Ryan Adamson†, and Christian Engelmann\*

\*Computer Science and Mathematics Division

†National Center for Computational Sciences

Oak Ridge National Laboratory

Oak Ridge, TN, USA

Email: {parkbh, hukerikarsr, adamsonrm, engelmannnc}@ornl.gov

**Abstract**—Today’s high-performance computing (HPC) systems are heavily instrumented, generating logs containing information about abnormal events, such as critical conditions, faults, errors and failures, system resource utilization, and about the resource usage of user applications. These logs, once fully analyzed and correlated, can produce detailed information about the system health, root causes of failures, and analyze an application’s interactions with the system, providing valuable insights to domain scientists and system administrators. However, processing HPC logs requires a deep understanding of hardware and software components at multiple layers of the system stack. Moreover, most log data is unstructured and voluminous, making it more difficult for system users and administrators to manually inspect the data. With rapid increases in the scale and complexity of HPC systems, log data processing is becoming a big data challenge. This paper introduces a HPC log data analytics framework that is based on a distributed NoSQL database technology, which provides scalability and high availability, and the Apache Spark framework for rapid in-memory processing of the log data. The analytics framework enables the extraction of a range of information about the system so that system administrators and end users alike can obtain necessary insights for their specific needs. We describe our experience with using this framework to glean insights from the log data about system behavior from the Titan supercomputer at the Oak Ridge National Laboratory.

## I. INTRODUCTION

Log data is essential for understanding the behavior of high-performance computing (HPC) systems by recording their usage and troubleshooting system faults. Today’s HPC systems are heavily instrumented at every layer for health monitoring by collecting with performance counters and resource usage data. Most components also report information about abnormal events, such as critical conditions, faults, errors and failures.

This manuscript has been authored by UT-Battelle, LLC under Contract No. DE-AC05-00OR22725 with the U.S. Department of Energy. The United States Government retains and the publisher, by accepting the article for publication, acknowledges that the United States Government retains a non-exclusive, paid-up, irrevocable, worldwide license to publish or reproduce the published form of this manuscript, or allow others to do so, for United States Government purposes. The Department of Energy will provide public access to these results of federally sponsored research in accordance with the DOE Public Access Plan (<http://energy.gov/downloads/doe-public-access-plan>).

This system activity and event information is logged for monitoring and analysis. Large-scale HPC installations produce various types of log data. For example, job logs maintain a history of application runs, the allocated resources, their sizes, user information, and exit statuses, i.e., successful vs. failed. Reliability, availability and serviceability (RAS) system logs derive data from various hardware and software sensors, such as temperature sensors, memory errors and processor utilization. Network systems collect data about network link bandwidth, congestion and routing and link faults. Input/output (I/O) and storage systems produce logs that record performance characteristics as well as data about degradations and errors detected.

HPC log data, when thoroughly investigated both in spatial and temporal dimensions, can be used to detect occurrences of failures and understand their root causes, identify persistent temporal and spatial patterns of failures, track error propagation, evaluate system reliability characteristics, and even analyze contention for shared resources in the system. However, HPC log data is derived from multiple monitoring frameworks and sensors and is inherently unstructured. Most log entries are not set up to be understood easily by humans, with some entries consisting of numeric values while others include cryptic text, hexadecimal codes, or error codes. The analysis of this data and finding correlations faces two main difficulties: first, the volume of RAS logs makes the manual inspection difficult; and second, the unstructured nature and idiosyncratic properties of log data produced by each subsystem log adds another dimension of difficulty in identifying implicit correlation among the events recorded. Consequently, the usage of log data is, in practice, largely limited to detection of mere occurrences of known text patterns that are already known to be associated with certain types of events.

As the scale and complexity of HPC systems continues to grow, the storage, retrieval, and comprehensive analysis of the log data is a significant challenge. In future extreme scale HPC systems the massive volume of monitoring and log data makes manual inspection and analysis impractical, and therefore poses a data analytics challenge. To address this

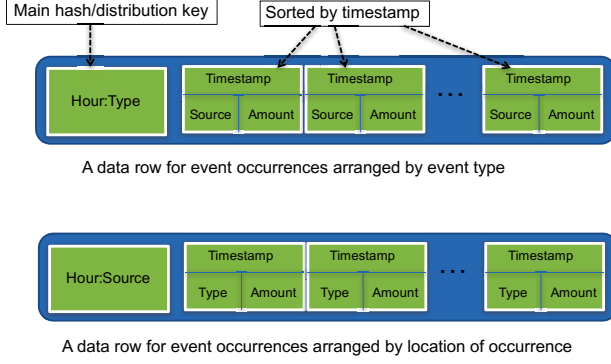


Fig. 1. Schemas for event occurrences: event schema ordered by time of occurrence (Top) and by location of occurrence (Bottom)

challenge, scalable methods for processing log and monitoring data are needed. This will require storing the enormous data sets in flexible schemas based on scalable and highly available database technologies, which can support large-scale analytics with low latencies, and high performance distributed data processing frameworks to support batch, real-time, and advanced analytics on the system data.

In this paper, we introduce a scalable HPC system data analytics framework, which is designed to provide system log data analysis capabilities to a wide range of researchers and engineers including system administrators, system researchers, and end users. The framework leverages Cassandra, a NoSQL distributed database to realize a scalable and fast-response backend for high throughput read/write operations, the Apache Spark for supporting rapid analysis on the voluminous system data. The framework provides a web-based graphical, interactive frontend interface that enables users to track system activity and performance and visualize the data. Using the framework, users can navigate spatio-temporal event space that overlaps with particular system events, faults, application runs, and resource usage to monitor the system, extract statistical features, and identify persistent behavioral patterns. End users can also visually inspect trends among the system events and contention on shared resources that occur during the run of their applications. Through such analysis, the users may find sources of performance anomalies and gain deeper insights into the impact of various system behaviors on application performance.

The rest of the document is organized as follows: Section II presents the data model and the design considerations that influenced the architecture of our framework. Section III details the architecture of our framework and how it has been adapted to analyze data from the Titan supercomputer at the Oak Ridge Leadership Computing Facility (OLCF). Section IV surveys related works in HPC monitoring frameworks and the analysis of log data. Finally, Section V concludes the paper with a discussion on potential future directions.

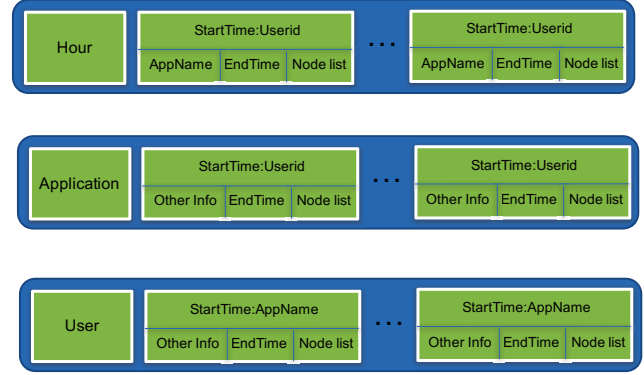


Fig. 2. Schemas for application runs: Application schema ordered by time of occurrence (Top), by name of application (Middle), and by users (Bottom)

## II. DATA MODEL

The monitoring infrastructure in supercomputing systems produces data streams from various sensors, which captures the resource and capacity utilization, power consumption, cooling systems, application performance, as well as various types of faults, errors and failures in the system. With the rapid increase in the complexity of supercomputing systems due to the use of millions of cores, complex memory hierarchies, and communication and file systems, massive amounts of monitoring data must be handled and analyzed in order to understand the characteristics of these systems and correlations between the various system measurements. The analysis of system monitoring data requires capturing relevant sensor data and system events, storing them in databases, developing analytic models to understand the spatial and temporal features of the data, or correlation between the various data streams, and providing tools capable of visualizing these system characteristics, or even building predictive models to improve the system operation. With the explosion in the monitoring data, this rapidly becomes a big data challenge. Therefore, to handle the massive amounts of system monitoring data and to support capabilities for more rigorous forms of user defined analytics, we adopt storage solutions designed to handle large amounts of data, and an in-memory data processing framework.

### A. Design Considerations

An implementation of an HPC system log analytics framework should start with extracting, transforming, and loading (ETL) of log data into a manageable database that can serve a flexible and rich set of queries over large amount of data. Due to the variety and volume of the data, we considered flexibility and fast performance to be two key design objectives of the framework. For an analytics framework to be successfully used for current and emerging system architectures, we placed emphasis on the following design considerations for the backend data model:

- **Scalability:** The framework needs to store historical log data as well as future events from the monitoring

frameworks. The data model should be scalable to accommodate an ever increasing volume of data.

- **Low latency:** The framework work also needs to serve interactive analytics that require near-real time query responses for timely visual updates. The backend data model should operate with minimal latency.
- **Flexibility:** A single data representation, or schema, for the various types of events from different system components is not feasible. The data model should offer flexible mechanism to add new event types and modify existing schemas to accommodate changes in system configuration, software updates, etc.
- **Time series friendly:** The most common type of log analytics that are of interest to HPC practitioners are expected to be based on time series data, which provide insights about the system's behavior over a user specified window of time.

We believe that these features will enable users to identify patterns among the event occurrences over time and explain the abnormal behavior of systems and the impact on applications. The foundation of the analytics framework on such a data model will support a variety of statistical or data mining techniques, such as association rules [1], decision trees [2], cross correlation [3], Bayesian network [4], etc., to be applied to the system log data.

For supporting a broad range of analytics, the retention of the raw data in a semi-structured format will be greatly beneficial. However, we found the conventional relational databases (RDBMS) do not satisfy our requirements. First, a schema of a relational database, once created, is very difficult to modify, whereas the format of HPC logs tend to change periodically. Second, due to its support for the atomicity, consistency, isolation, and durability (ACID) properties and two-phase commit protocols, it does not scale. After investigating various database technologies, we found the Apache Cassandra [5] to be most suitable for building the backend data model for our design of log analytics framework. Cassandra, based on Amazon's Dynamo and Google's BigTable, is a column-oriented distributed database offering highly available (HA) services with no single point of failure. Cassandra, a hashing-based distributed database system, stores data in *tables*. A data unit of a table, also known as *partition*, is associated with a hash key and mapped to one or more nodes of a Cassandra cluster. A partition can be considered as a data row that can contain multiple column families, where each family can be of a different format. Cassandra's performance, resiliency, and scalability come from its master-less ring design, which unlike a legacy master-slave architecture gives an identical role to each node. With a replication option that is implemented on commodity hardware, Cassandra offers a fault tolerant data service. Also with its column oriented features, Cassandra is naturally suitable for handling data in sequence, regardless of data sizes. When data is written to Cassandra, each data record is sorted and written sequentially to disk. When a database is queried, data is retrieved by row key and range within a row,

which guarantees a fast and efficient search.

## B. Data Model Design

Our data model is designed to initially study the operational behavior of the Titan supercomputer hosted by the Oak Ridge National Laboratory. The framework is designed to study Titan's system logs collected from console, application and network logs, which contain timestamped entries of critical system events. The data model is designed to capture various system events including, machine check exceptions, memory errors, GPU failures, GPU memory errors, Lustre file system errors, data virtualization service errors, network errors, application aborts, kernel panics, etc.

We have created a total of eight tables to model system information, the types of event we monitor, occurrences of events, and application runs. The partitions for events are designed to disperse overheads in both reading and writing data evenly over to the cluster nodes. Fig 1 shows how a partition is mapped to one of the four nodes by its hash key of hour and type combination.

- *nodeinfos*
- *eventtypes*
- *eventsynopsis*
- *event\_by\_time*
- *event\_by\_location*
- *application\_by\_time*
- *application\_by\_user*
- *application\_by\_location*

The *nodeinfos* contains information about the system including the position of a rack (or cabinet) in terms of row and column number, the position of a compute node in terms of rack, chassis, blade, and module number, network and routing information, etc. Each node in the Titan system consists of a AMD CPU and a NVIDIA GPU. Each CPU is a 16-core AMD Opteron 6274 processor with 32 GB of DDR3 memory and each GPU is a NVIDIA K20X Kepler architecture-based GPU with 6 GB of GDDR5 memory. The system uses Cray Gemini routers, which are shared between a pair of nodes. Each blade/slot Titan supercomputer consists of four nodes. Each cage has eight such blades and a cabinet contains three such cages. The complete system consists of 200 cabinets that are organized in a grid of 25 rows and 8 columns. The *nodeinfo* enables spatial correlation and analysis of events in the system.

The two tables *event\_by\_time* and *event\_by\_location* store system event information from two perspectives, time and location to facilitate spatio-temporal analysis. An event in our data model is defined as occurrence(s) of a certain type reported at a particular timestamp. An event is also associated with the location (or the source component) where it is reported. The two tables illustrate these dual representations of an event as illustrated in Fig 1. The first table structure associates an event with its type and the hour of its occurrence; all events of a certain type generated at a certain hour are stored in the same partition. In contrast, the second table structure associates an event with hour and location; all events, regardless of their type, generated at a certain hour for the

same component, are stored in the same partition. Note that each partition stores events sorted by their timestamps, which is a time series representation of events that is one hour long. This facilitates to support a spatio-temporal query.

For data about user application runs, we added another dimension: users. More specifically, three tables to represent user application runs from perspectives of time, application, and user (see Fig 2). Readers can find this as a set of denormalized views on application runs. Note however, although all application runs in each partition type are depicted the same, in fact, each application run may include columns unique to it. For example, a column named as *Other Info* may include multiple sub-columns to represent different information.

### III. ARCHITECTURE OF THE LOG ANALYTICS FRAMEWORK

The layout of the overall architecture, illustrated in Fig 3, consists of three main components: a web-based frontend, the data analytics server, and the backend distributed database. The frontend, which consists of a client-side application, adopts a web interface allowing users to create queries for the analysis of the data as well as for visual inspection of log data and application runs in both spatial and temporal domain. The analytics server translates data query requests received from the frontend and relays them to the backend database server in the form of Cassandra Query Language (CQL) queries. The query results from the backend are returned to the frontend, or as intermediate data for further processing. The backend distributed NoSQL database stores and manages Titan system logs and user application logs. The backend server communicates with the analytics layer through a RESTful interface. Query results are sent in JSON object format to avoid data format conversion at the frontend. This framework is currently being deployed at the ORNL's Compute and Data Environment for Science (CADES), which provides flexible computing, data storage and analytics infrastructure.

#### A. Analytic Server and Backend Database

The analytics server consists of a web server, a query processing engine, and a big data processing engine. The user queries are received by the web server, translated by the query engine, and either forwarded to the backend database, or the big data processing unit depending on the type of a user query. Simple queries are directly handled by the query engine, and complex queries are passed to the big data processing unit. The big data processing unit initiates a Spark session over the Spark cluster that reside on the same nodes with Cassandra.

Since the analytic framework intends to serve numerous users, who may require long-lived connections and may expect delayed responses from the server for non-trivial analytics, we chose the Tornado framework [6], which provides a web server and asynchronous networking library. Tornado supports non-blocking I/O, which makes it suitable for *long-polling* and *WebSockets* to implement the long-living connections from the frontend web-based client.

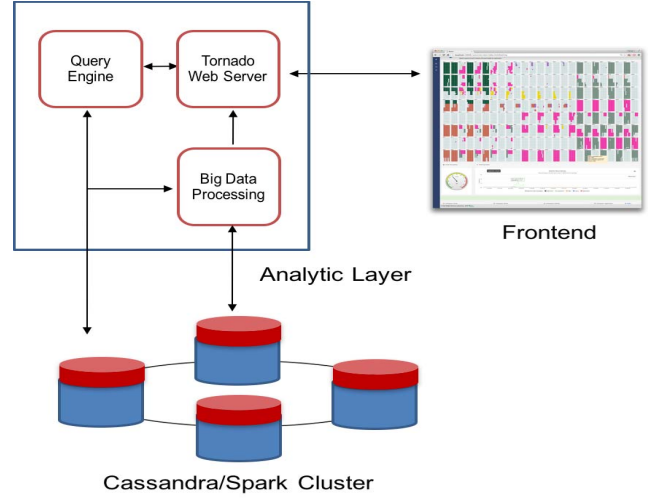


Fig. 3. Overall architecture of the Log Analytics Framework consisting of the Cassandra distributed NoSQL database and the Apache Spark in-memory data processing engine

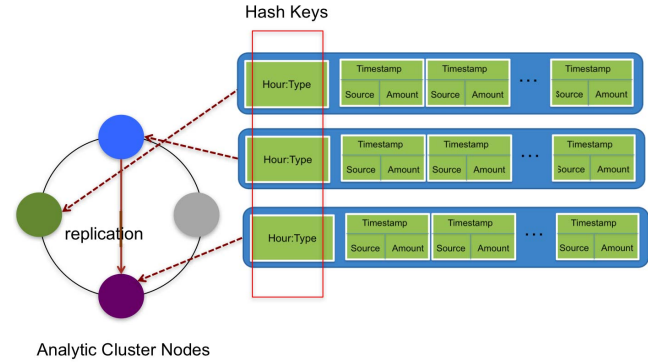


Fig. 4. Event partitions mapped to Cassandra nodes by hour and event types

The Cassandra backend database server and the Apache Spark cluster are installed over the same 32 virtual machine (VM) instances in CADES, that is, a pair of a Spark worker node and a Cassandra node runs together in each of the 32 VMs. We selected this configuration to maximize data locality for the computation performed by the analytic algorithms of the big data processing unit. As described in Section II, a *partition* of a table is defined by a combination of hour, user, application, location, and event type representing the data from a specific view (this will be defined as a *context* below). Each table is distributed over the entire cluster retaining time ordered data entries within each partition. The big data processing unit consists of a set of Spark computations that perform MapReduce operations over time ordered data spread across the cluster by a context. By associating local partitions with the same local Spark worker, the big data processing unit performs analytics efficiently. Fig 4 illustrates an example of partitions for event occurrences that are mapped to nodes on the basis of the event hour and event type.



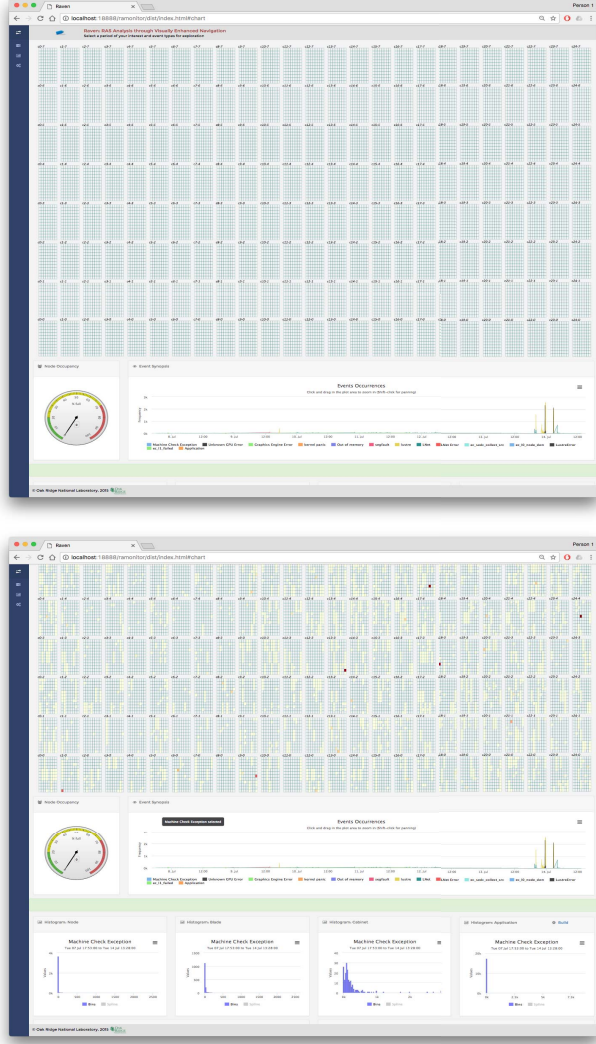


Fig. 5. The physical system map and the temporal map of the frontend (Top); Distribution of an event type over a selected period as a heat map on the physical system map and event histograms (Bottom).

### B. Frontend: Client Module

The frontend provides a window to the system log data. Users interact with the frontend to inspect the system status or perform analytics on log data. Every interaction with the frontend is translated into a query in Javascript Object Notation (JSON) format and delivered to the analytic server. The current frontend provides visualization of events and application runs in both spatial and temporal dimension on physical system map and time interval map. The visualization is implemented using D3 [7] package and HTML5 canvas.

Users interact with the framework by creating a *context*. A context is selected on the basis of event type, application, location, user, time period, or a combination of these, over which the system status is defined and examined. By selecting a context, important insights about the system status can be



Fig. 6. Event occurrences (Top) and Application placement (Bottom) rendered on the Physical System Map

extracted. The selection of an appropriate context also helps in identifying the root cause of failure events. When a context is created the appropriate query is passed to the data analytic server to retrieve data. The frontend allows users to choose desired contexts and results by interacting with:

- The physical system map
- The temporal map
- The event types map
- The user/application map
- The tabular map of raw log entries

While the physical system map shows the spatial placement of racks (or cabinets) and the individual nodes within each rack, the temporal map shows occurrences of events over a time interval. Fig 5-(Top) shows the physical system map and the temporal map. Event occurrences, or application displacements, are displayed on the physical system map. Using the event type map and the temporal map, users can select an event type of interest at a particular time. The occurrences of the selected event type at the specified timestamp are shown on the compute nodes where they occurred in the physical system map. Likewise, displacements of all applications that were running at the time of selection, once users select using the user/application map, are shown on the nodes in the physical system map. Fig 6 shows Lustre error occurrences on each compute node (Top) and the placement of user applications (Bottom) at the specified timestamp. Users can also select a hardware component such as compute node in the physical system map. Thus, the physical system map, the temporal map, the event types map, and the user/application map are essentially interactive visualization components that allow discovery of correlations among events, locations, and applications by tracing occurrences and progression of events.

The temporal map represents a selected time interval. Users

can repeatedly select sub-intervals of interest for narrowed investigations. With a selected interval, users can extract basic statistics about event occurrences. First, users can create a heat map representation of the occurrences of an event type within the interval on the physical system map, which illustrates whether the event occurrences were unusually higher (or lower) in some parts of the system compared to the rest of the parts. In addition, users can also get distributions of the event occurrences over cabinets, blades, nodes, and applications. These two different types of view (heat map and distributions) offer complementary insights on normal or abnormal occurrences of a certain event type observed during a selected period. Fig 5-(Bottom) shows that Machine Check Exception (MCE) errors occurred abnormally high in some compute nodes over a selected time period.

### C. Big Data Processing using the Frontend

The big data processing unit intends to serve a wide range of users for intensive analytic processing. We are currently developing log data analytic application program interfaces (APIs) through which users can connect to the analytic server from their chosen applications. The frontend also offers a set of basic analytics capabilities utilizing big data processing unit.

First, the heat map representation and various distributions of event occurrences over a selected time interval, which are mentioned above, are computed by the big data processing. Second, the investigation of correlation between two event occurrences within a selected time interval, which can provide a causal relationship between the two, is also processed by the big data processing unit. Fig 7-(Top) shows the transfer entropy plot of two events measured within a selected time window.

Also, basic text analytics are supported by the big data processing unit. Identification of *important* keywords (either letters or alphanumeric values) from raw system logs often helps understanding the system status given a massive number of system events logged. Examples include events from the filesystem, the network subsystem, etc. For example, Lustre log message contains descriptions regarding status of hardware, an I/O transaction, the peers of the log generator, etc. These information are written in texts, hexadecimal numbers, or special characters. Once properly filtered, each Lustre event message can be transformed into a set of words that represents the event occurrence as a point in a metric space. Such transformations typically involve *word counts* and/or term frequency-inverse document frequency (TF-IDF) of log messages. Note here a Lustre message is treated as a document from a conventional text analysis point of view. The temporal event view in Fig 7-(Bottom) shows a period when tens of thousands Lustre error messages were generated. As shown in the map, it was a system wide event that lasted several minutes afflicting most of compute nodes and applications running therein. In many cases, the root causes of such a system-wide event are abnormal behaviors of either hardware or a system software component of which negative impacts propagate over to the entire system. Although it may be a single source

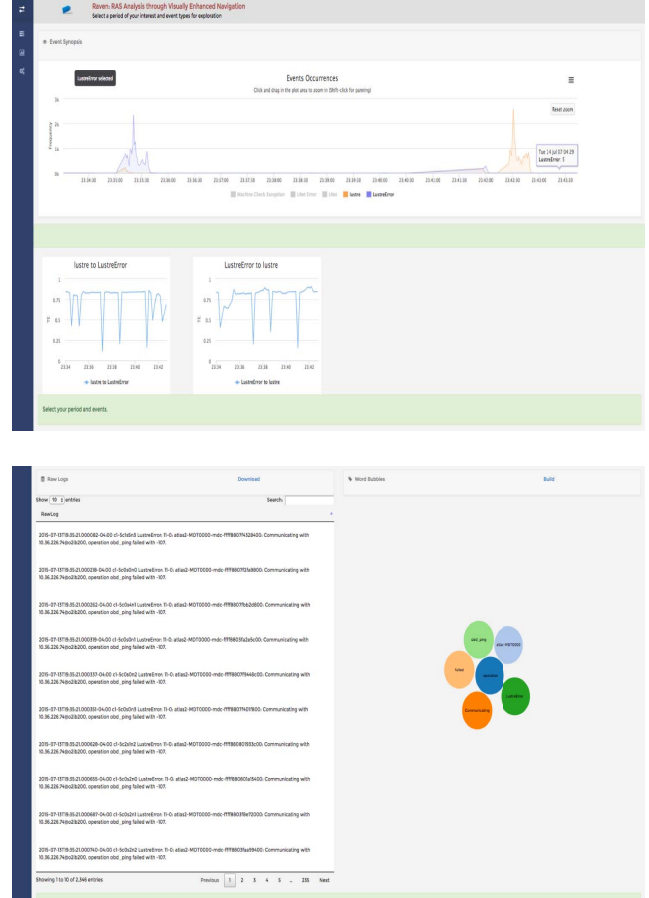


Fig. 7. Transfer Entropy plot of two event types measured within a selected time interval (Top). Raw log entries shown in the tabular map and importance words from logs illustrated as bubbles (Bottom).

problem, it requires to sift through a large volume of Lustre event logs to identify the problem components. We found that a simple word counts, which is rapidly executed by Spark, can locate the source of the problem. Fig 7-(Bottom) shows word bubbles as the result of text analysis on raw Lustre event logs, which illustrates an object storage target is not responding.

### D. Data Ingestion

The log analytics framework is designed to ingest new event data in two different modes: batch import and real-time streaming. The batch import is a traditional ETL procedure that involves 1) collocation of all data, 2) parsing the data in search for known patterns for each event type (typically defined as regular expressions), and 3) batch upload into the backend database. The batch import is used when a new event type is identified and all occurrences in the historical data must be collected. Since such an update may require huge computational overheads, the analytic framework implements parsing and uploading using Apache Spark.

The real-time streaming mode, which is currently being developed through a collaboration with the high-performance

computing operation group of the Oak Ridge Leadership Computing Facility (OLCF), intends to facilitate online analytics such as real time failure detection by monitoring recent event streams. The OLCF is developing event producers that not only parse real-time streams from log sources but also publish each event occurrence from the streams. Each event occurrence is published to an Apache Kafka message bus that is available to consumers subscribing to the corresponding topic. For example, event logs of a Lustre filesystem are generated at multiple places: the servers (OSSes, MDSes, and MGS), and the clients at each compute node. In addition, lower level components (e.g. disk controllers) also generate separate logs.

The OLCF has deployed Kafka on top of OpenShift Origin, which is a scalable container-based framework for scheduling applications. This deployment method allows elastic scale-out of Kafka nodes and backend databases to accommodate large upswings in datatype growth. As new supercomputing resources and high performance filesystems are brought online, additional OpenShift pods will be deployed as necessary to handle additional log event types. Conversely, as analytics frameworks grow in size and complexity, the Kafka install will be scaled out to meet the demand of any consumers.

To receive event streams from the Kafka instance, the analytic framework places a subscriber that delivers event messages to Spark streaming module that in turn converts and places all event occurrences into the right partitions. Event occurrences of the same type and same location are coalesced into a single event if they are timestamped the same. For this, the time window of the Spark streaming is set to one second. This real-time upload component will be further extended to include various online analytic modules.

#### IV. RELATED WORK

Various monitoring frameworks are used in large-scale computing systems for understanding the use of system resources by applications, the impact of competition for shared resources and for the discovery of abnormal system conditions in the system. Tools such as Ganglia [8] and Nagios [9] are widely used for in HPC cluster and grid systems as well as in enterprise clusters. OVIS [10] provides a suite of monitoring and analysis tools for HPC systems that provides finer grained monitoring to enable understanding platform resource utilization characteristics of applications.

Several studies have sought to analyze failures in large-scale systems to characterize the reliability of the system. These studies attempt to characterize the root causes of single-node failures as well as system-wide failures from manual reports and system event logs [11]. These studies perform post-mortem analysis on the system logs to extract the statistical properties of system errors and failures [12] [13]. For the analysis of logs of large-scale systems, certain approaches apply filtering, followed by extraction and categorization of error events [14] [15] [16]. Other analyses use approaches such as time coalescing [17]. Some studies have focused on analysis of failure characteristics of specific subsystems or system components in HPC systems, such as disks [18],

DRAM memory [18] [19] [20], graphical processing units (GPU) [21]. These studies sanitize the system logs using manual failure reports, or extract specific events of interest, to compute the relative failure frequencies for various root causes and their mean and standard deviation in contrast to our framework, which mines for insights from the unstructured raw data. Our approach is designed to handle massive amounts of heterogeneous monitoring and log data, which will be typical in future extreme-scale systems with complex hardware and software architectures.

Based on the observation of characteristics of failure events and correlations between the events, models for failure prediction have been proposed [22] [23]. These prediction algorithms leverage the spatial and temporal correlation between historical failures, or trends of non-fatal events preceding failures to design remedial actions in the system's configuration, scheduling algorithms to mitigate the adverse impacts of failure events.

#### V. CONCLUSION

With the ever-growing scale and complexity of high performance computing (HPC) systems, characterizing system behavior has become a significant challenge. The systems produce and log vast amounts of unstructured multi-dimensional data collected using a variety of monitoring tools. The tools and methods available today for processing this log data lack advanced data analytics capabilities, which makes it difficult to diagnose and completely understand the impact of system performance variations, fault and error events in the system on application performance. To handle the massive amounts of system log data from a diverse set of monitoring frameworks and rapidly identify problems and variations in system behavior, it is essential to have scalable tools to store and analyze the data.

In this paper, we introduced a scalable HPC log data analytics framework based on a distributed data and computation model. The framework defines a time-series oriented data model for HPC log data. We leverage big data frameworks, including Cassandra, a highly scalable, high-performance column-oriented NoSQL distributed database, and Apache Spark, a real-time distributed in-memory analytics engine. We presented a data model designed to facilitate log data analytics for system administrators and researchers as well as end users who are often oblivious to the impact of variations and fault events on their application jobs.

Our log analytic framework has been tested with Titan supercomputer at the Oak Ridge Leadership Computing Facility's (OLCF). Although the framework is still evolving, with new analytics modules being currently developed, the preliminary assessment shows that the framework can provide deeper insights about the root causes of system faults, and abnormal behaviors of user applications. It also enables statistical analysis of event occurrences and their correlations on a spatial and temporal basis. These capabilities will be valuable when deploying a new HPC system in the pre-production phase, as well as during operational lifetime for fine tuning the system.

While our existing framework improves upon the state-of-the-art in HPC log data processing, there is much room to improve. As future work, we are planning several enhancements and improvements to the framework. First, new and composite event types will need to be defined for capturing the complete status of the system. This will involve event mining techniques rather than text pattern matching. Second, the framework will need to develop application profiles in terms of event occurred during its runs. This will help understand correlations between application runtime characteristics and variations observed in the system on account of faults and errors. Finally, the framework will need to support advanced statistical techniques, incorporate machine learning algorithms, and graph analytics for more comprehensive investigation of log and monitoring data.

#### ACKNOWLEDGMENT

This material is based upon work supported by the U.S. Department of Energy, Office of Science, Office of Advanced Scientific Computing Research, program manager Lucy Nowell, under contract number DE-AC05-00OR22725. This work was supported by the Compute and Data Environment for Science (CADES) facility and the Oak Ridge Leadership Computing Facility at the Oak Ridge National Laboratory, which is managed by UT Battelle, LLC for the U.S. DOE (under the contract No. DE-AC05-00OR22725).

#### REFERENCES

- [1] G. Piatetsky-Shapiro, "Discovery, analysis, and presentation of strong rules," in *Knowledge Discovery in Databases*, 1991.
- [2] J. R. Quinlan, "Induction of decision trees," *Machine Learning*, vol. 1, no. 1, pp. 81–106, March 1986.
- [3] N. Bansal, A. Blum, and S. Chawla, "Correlation clustering," *Machine Learning*, vol. 56, no. 1, pp. 89–113, 2004.
- [4] D. Heckerman, "Bayesian networks for data mining," *Data Mining and Knowledge Discovery*, vol. 1, no. 1, pp. 79–119, 1997.
- [5] Apache Software Foundation. (2008) Apache cassandra. [Online]. Available: <http://cassandra.apache.org>
- [6] B. Taylor *et al.* (2009) Tornado web server. [Online]. Available: <http://www.tornadoweb.org/>
- [7] M. Bostock *et al.* (2011) D3.js. [Online]. Available: <https://d3js.org>
- [8] M. L. Massie, B. N. Chun, and D. E. Culler, "The ganglia distributed monitoring system: design, implementation, and experience," *Parallel Computing*, vol. 30, no. 7, pp. 817 – 840, 2004.
- [9] W. Barth, *Nagios: System and Network Monitoring*, 2nd ed. San Francisco, CA, USA: No Starch Press, 2008.
- [10] J. M. Brandt, A. C. Gentile, D. J. Hale, and P. P. Pebay, "Ovis: a tool for intelligent, real-time monitoring of computational clusters," in *Proceedings 20th IEEE International Parallel Distributed Processing Symposium*, April 2006, p. 8.
- [11] C. D. Martino, Z. Kalbarczyk, R. K. Iyer, F. Baccanico, J. Fullop, and W. Kramer, "Lessons learned from the analysis of system failures at petascale: The case of blue waters," in *2014 44th Annual IEEE/IFIP International Conference on Dependable Systems and Networks*, June 2014, pp. 610–621.
- [12] R. K. Sahoo, M. S. Squillante, A. Sivasubramaniam, and Y. Zhang, "Failure data analysis of a large-scale heterogeneous server environment," in *International Conference on Dependable Systems and Networks*, 2004, June 2004, pp. 772–781.
- [13] B. Schroeder and G. Gibson, "A large-scale study of failures in high-performance computing systems," *IEEE Transactions on Dependable and Secure Computing*, vol. 7, no. 4, pp. 337–350, Oct 2010.
- [14] Y. Liang, Y. Zhang, A. Sivasubramaniam, R. K. Sahoo, J. Moreira, and M. Gupta, "Filtering failure logs for a bluegene/l prototype," in *2005 International Conference on Dependable Systems and Networks (DSN'05)*, June 2005, pp. 476–485.
- [15] A. Oliner and J. Stearley, "What supercomputers say: A study of five system logs," in *37th Annual IEEE/IFIP International Conference on Dependable Systems and Networks (DSN'07)*, June 2007, pp. 575–584.
- [16] A. Pecchia, D. Cotroneo, Z. Kalbarczyk, and R. K. Iyer, "Improving log-based field failure data analysis of multi-node computing systems," in *2011 IEEE/IFIP 41st International Conference on Dependable Systems Networks (DSN)*, June 2011, pp. 97–108.
- [17] C. D. Martino, M. Cinque, and D. Cotroneo, "Assessing time coalescence techniques for the analysis of supercomputer logs," in *IEEE/IFIP International Conference on Dependable Systems and Networks (DSN 2012)*, June 2012, pp. 1–12.
- [18] B. Schroeder, E. Pinheiro, and W.-D. Weber, "Dram errors in the wild: A large-scale field study," in *Proceedings of the Eleventh International Joint Conference on Measurement and Modeling of Computer Systems*, ser. SIGMETRICS '09. ACM, 2009, pp. 193–204.
- [19] A. A. Hwang, I. A. Stefanovici, and B. Schroeder, "Cosmic rays don't strike twice: Understanding the nature of dram errors and the implications for system design," *SIGPLAN Not.*, vol. 47, no. 4, pp. 111–122, Mar. 2012.
- [20] V. Sridharan, J. Stearley, N. DeBardeleben, S. Blanchard, and S. Gurumurthi, "Feng shui of supercomputer memory: Positional effects in dram and sram faults," in *Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis*, ser. SC '13. ACM, 2013, pp. 22:1–22:11.
- [21] D. Tiwari, S. Gupta, G. Gallano, J. Rogers, and D. Maxwell, "Reliability lessons learned from gpu experience with the titan supercomputer at oak ridge leadership computing facility," in *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, ser. SC '15. ACM, 2015, pp. 38:1–38:12.
- [22] Y. Liang, Y. Zhang, A. Sivasubramaniam, M. Jette, and R. Sahoo, "Bluegene/l failure analysis and prediction models," in *International Conference on Dependable Systems and Networks (DSN'06)*, June 2006, pp. 425–434.
- [23] A. Gainaru, F. Cappello, M. Snir, and W. Kramer, "Fault prediction under the microscope: A closer look into hpc systems," in *High Performance Computing, Networking, Storage and Analysis (SC), 2012 International Conference for*, Nov 2012, pp. 1–11.